

CSE 202: Fundamentals of Database Systems

Winter 2019

Assignment 2

Due Date: 25-02-2019 (Three Weeks' time: No extension will be allowed)

Instructions:

- Write the programs in **Java** only.
- Only one submission per group.
- The naming convention for the files: filename_rollnumber.extension ,like program1_2016001_2016002.java.
- **The assignment can be done individually or in groups of maximum 2 students.**
- **Compress all the input/output files along with the programs as tar.gz. Your submission (.tar.gz) must contain a text file with your group details (name and roll number of student). Your submission must contain files only with no folder hierarchy created.**

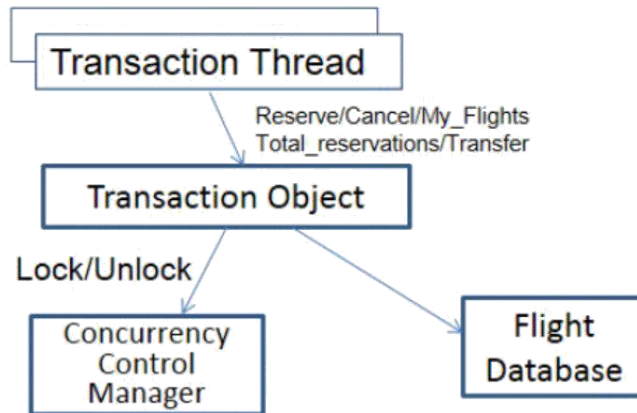
In this assignment, you have to implement a concurrency control algorithm for an airline database. The system consists of a transaction threads, concurrency control manager and an airline database. The database has a set of transactions and data structures containing flight information (passenger lists for all flights). Each transaction thread (TT) should have the following structure:

Repeat

Select a transaction type randomly
Select object (flight and passenger id) for transaction randomly
Invoke transaction

The possible transactions are:

- **Reserve(F, i)**: reserve a seat for passenger with id i on flight F, where $i > 0$.
- **Cancel(F, i)**: cancel reservation for passenger with id i from flight F.
- **My_Flights(id)**: returns the set of flights on which passenger i has a reservation.
- **Total_Reservations()**: returns the sum total of all reservations on all flights.
- **Transfer(F1,F2,i)**: transfer passenger i from flight F1 to F2. This transaction should have no impact if the passenger is not found in F1 or there is no room in F2.



You must choose appropriate data structures to implement the flight database. However, these data structures must contain no synchronization code. The Concurrency Control Manager (CCM) must implement the mechanism to control access to the data structures. At the minimum, it should include locking and unlocking operations. The Transaction object must implement the operations Reserve, Cancel, My_Flights, Total reservations and Transfer operations (described above). Each of these operation must invoke the necessary lock and unlock operations and operations on the data structures.

The goal is to study performance improvements based on the granularity of locking. The first version of the program must implement only serial schedules (lock the entire database). You must then compare the performance to your second version which does two-phase locking at a more fine-grained level. You must exploit semantics of the operations to allow as much concurrency as possible.

The performance must be measured in terms of transaction throughput (number of transactions completed over a specific interval of time). You must have a sleep statement inside the function calls of the data structures to simulate time required to access the database. You must plot the following data for each version: (a) Impact of number of TT threads on throughput, (b) Impact of the transaction mix on throughput. You should have information for around 5 flights. Each flight should have a different bound of the number of seats available. The number of TT threads should be increased to a sufficient number to see a trend the performance graphs.